

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

1 K-Means Clustering

Load Data

```
In [2]: df = pd.read_csv('/Users/chenjiquing/Public/2022_Spring_term/QBS108 Applied ML/2022 file/Homework 2/tcga_pancancer/pancancer.csv')
df = df.drop('Unnamed: 0',axis=1)
df
```

Out[2]:

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
0	-25.875320	56.783945	-58.522554	-23.657779	59.226184	32.978484	-55.835789	3.992981	-16.705564	-14.854115	-16.739289	-12.236637	0.631551	-0.895200	31.014503
1	-78.976524	-62.312761	33.852422	-2.165579	-5.288384	0.446068	24.361695	-44.745726	17.395310	-22.758601	13.726631	11.176496	-6.800069	9.362851	-7.950115
2	87.151357	-23.404469	-10.054268	63.831495	-8.287112	-42.402161	36.263516	-54.548902	33.558731	2.175367	-26.291564	-12.975848	24.410122	-28.263414	-17.609067
3	-43.224891	-24.016296	-4.744443	-66.180500	75.203354	-46.929467	37.587578	-20.590401	8.761443	-20.297368	4.353383	-2.354844	5.791677	1.503404	-1.664383
4	-63.589175	-66.331180	6.376695	25.804671	-35.920993	-20.097241	4.684812	-1.233395	-26.762385	18.346569	11.932363	0.081588	-3.752243	4.445100	-0.940169
...
660	-77.968110	-64.215899	53.033778	16.772297	7.374632	51.174647	-22.336068	-17.173816	-16.990365	9.496480	19.249984	-17.164237	9.619327	-10.789798	8.587151
661	1.411024	33.721905	-97.514669	-35.977829	-32.405695	54.035778	27.624459	1.633711	-30.978002	-22.327608	15.138285	-20.205383	-6.395287	5.546136	2.314934
662	-69.761387	-80.963905	40.077150	7.640976	-18.890106	18.978991	-2.254887	-34.397698	2.597095	-6.976562	-11.659815	-6.662376	9.278010	-4.117283	-30.055490
663	-24.481133	61.988138	-65.444488	34.473801	-6.558216	54.685834	42.522086	0.694724	43.626111	22.321826	-20.691297	44.923379	-14.805007	-20.591455	-25.827407
664	5.547498	47.154601	-49.546932	1.701638	-56.533622	-96.443116	34.725324	69.097112	65.273405	45.460026	32.333668	-22.421023	0.659979	27.597299	13.250129

665 rows x 15 columns

```
In [3]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 665 entries, 0 to 664
Data columns (total 15 columns):
#   Column  Non-Null Count  Dtype  
---  -
0    0        665 non-null    float64
1    1        665 non-null    float64
2    2        665 non-null    float64
3    3        665 non-null    float64
4    4        665 non-null    float64
5    5        665 non-null    float64
6    6        665 non-null    float64
7    7        665 non-null    float64
8    8        665 non-null    float64
9    9        665 non-null    float64
10   10       665 non-null    float64
11   11       665 non-null    float64
12   12       665 non-null    float64
13   13       665 non-null    float64
14   14       665 non-null    float64
dtypes: float64(15)
memory usage: 78.1 KB
```

- No missing Values

```
In [4]: df.describe()
```

Out[4]:	0	1	2	3	4	5	6	7	8	9	10	11	12
count	6.650000e+02	6.650000e+02	6.650000e+02	6.650000e+02	6.650000e+02	6.650000e+02	6.650000e+02	6.650000e+02	6.650000e+02	6.650000e+02	6.650000e+02	6.650000e+02	6.650000e+02
mean	5.128730e-16	2.564365e-16	6.838306e-16	-7.265700e-16	-8.547883e-17	-5.128730e-16	2.564365e-16	-1.709577e-16	2.136971e-16	2.564365e-16	8.547883e-16	-8.547883e-17	3.846547e-16
std	8.126705e+01	6.536298e+01	5.221270e+01	3.631874e+01	3.486074e+01	3.342800e+01	2.608269e+01	2.340081e+01	2.243101e+01	1.965236e+01	1.914324e+01	1.779294e+01	1.708988e+01
min	-1.018639e+02	-9.465788e+01	-1.238769e+02	-1.209602e+02	-7.862494e+01	-9.644312e+01	-8.113492e+01	-7.474574e+01	-5.490130e+01	-5.870098e+01	-7.391167e+01	-5.091354e+01	-5.368153e+01
25%	-5.585776e+01	-5.007898e+01	-2.176080e+01	-2.296131e+01	-2.246344e+01	-2.361122e+01	-1.600493e+01	-1.637985e+01	-1.411962e+01	-1.336469e+01	-1.203067e+01	-1.022459e+01	-1.027628e+01
50%	-3.121861e+01	-1.626691e+01	9.400684e+00	-6.231582e+00	-3.753839e+00	1.149227e+00	-1.129484e+00	4.821054e-01	-6.971027e-01	-3.945739e-01	-1.057407e+00	-1.360354e-01	-1.799235e+00
75%	1.828140e+00	3.597135e+01	3.356166e+01	1.716525e+01	1.390316e+01	2.234912e+01	1.621502e+01	1.574662e+01	1.330756e+01	1.237379e+01	1.109708e+01	8.011535e+00	8.542092e+00
max	1.799845e+02	1.977116e+02	1.300737e+02	1.657368e+02	1.253445e+02	9.710720e+01	1.063313e+02	8.092908e+01	9.623726e+01	8.809292e+01	8.502497e+01	5.819931e+01	8.065588e+01

1-1 a. Choose a suitable type of plot and visualize the first two features of the data.


```
In [5]: # load module
import seaborn as sns
```

First feature

```
In [6]: sns.distplot(df['0'], hist = True, hist_kws = {"edgecolor": "black"})
plt.title("Histogram of First feature")
plt.xlabel("First Feature")
```

/Users/chenjiqing/anaconda3/envs/Class/lib/python3.7/site-packages/seaborn/distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histogram s).

Out[6]: Text(0.5, 0, 'First Feature')




Second feature

```
In [7]: sns.distplot(df['1'], hist = True, hist_kws = {"edgecolor": "black"})
plt.title("Histogram of Second feature")
plt.xlabel("Second Feature")
```

/Users/chenjiqing/anaconda3/envs/Class/lib/python3.7/site-packages/seaborn/distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histogram s).

Out[7]: Text(0.5, 0, 'Second Feature')



1-1 b. From your plot, how many clusters, k , do you recognize in the dataset?

- According to Histogram plot, I think 2 clusters in the dataset.

1-2 a. Using the k-Means algorithm, implement a clustering model. Train the clustering model using all 15 features from the data.

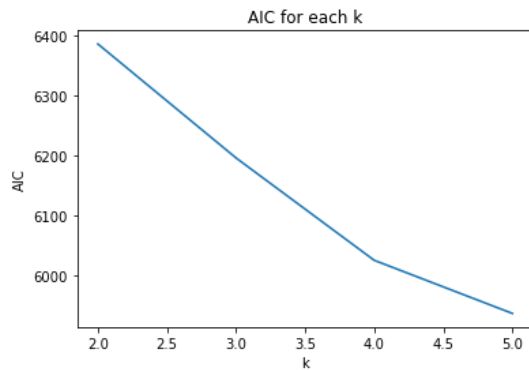
```
In [8]: from sklearn.cluster import KMeans

k_cluster = []
AICs = []
BICs = []
for k in range(2,6):
    kmeans = KMeans(n_clusters=k, random_state=3).fit(df)
    AIC = df.shape[0]*np.log(kmeans.inertia_/df.shape[0]) + 2*k
    BIC = df.shape[0]*np.log(kmeans.inertia_/df.shape[0]) + k*np.log(df.shape[0])
    k_cluster.append(k)
    AICs.append(AIC)
    BICs.append(BIC)
```

Plot AIC for each value of k .

```
In [9]: plt.plot(k_cluster, AICs)
plt.title("AIC for each k")
plt.xlabel("k")
plt.ylabel("AIC")
```

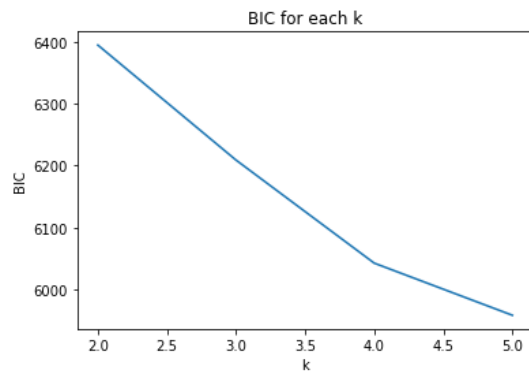
```
Out[9]: Text(0, 0.5, 'AIC')
```



Plot BIC for each value of k.

```
In [10]: plt.plot(k_cluster, BICs)
plt.title("BIC for each k")
plt.xlabel("k")
plt.ylabel("BIC")
```

Out[10]: Text(0, 0.5, 'BIC')



1-2 b. Which value of k is optimal?

- According to AIC and BIC plots, the optimal k is 5 (with lowest AIC and BIC values). The result doesn't meet my expectation in 1.1b

2 K-Nearest Neighbor Classification

Load Data

Training Set

```
In [11]: df_train = pd.read_csv('/Users/chenjiqing/Public/2022_Spring_term/QBS108 Applied ML/2022 file/Homework 2/wine/train.csv')
df_train = df_train.assign(good_quality = np.where(df_train['good_quality']!=False,0,1)) # Convert target into 1, 0
df_train
```

Out[11]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	good_quality
0	7.6	0.420	0.08	2.70	0.084	15.0	48.0	0.99680	3.21	0.59	10.0	0

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	good_quality
1	5.7	0.600	0.00	1.40	0.063	11.0	18.0	0.99191	3.45	0.56	12.2	1
2	7.5	0.610	0.26	1.90	0.073	24.0	88.0	0.99612	3.30	0.53	9.8	0
3	8.9	0.480	0.24	2.85	0.094	35.0	106.0	0.99820	3.10	0.53	9.2	0
4	7.0	0.620	0.18	1.50	0.062	7.0	50.0	0.99510	3.08	0.60	9.3	0
...
1018	7.4	0.635	0.10	2.40	0.080	16.0	33.0	0.99736	3.58	0.69	10.8	1
1019	10.3	0.440	0.50	4.50	0.107	5.0	13.0	0.99800	3.28	0.83	11.5	0
1020	6.6	0.520	0.08	2.40	0.070	13.0	26.0	0.99358	3.40	0.72	12.5	1
1021	5.1	0.510	0.18	2.10	0.042	16.0	101.0	0.99240	3.46	0.87	12.9	1
1022	11.6	0.410	0.54	1.50	0.095	22.0	41.0	0.99735	3.02	0.76	9.9	1

1023 rows × 12 columns

```
In [12]: df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1023 entries, 0 to 1022
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity         1023 non-null   float64
1   volatile acidity      1023 non-null   float64
2   citric acid           1023 non-null   float64
3   residual sugar        1023 non-null   float64
4   chlorides             1023 non-null   float64
5   free sulfur dioxide    1023 non-null   float64
6   total sulfur dioxide   1023 non-null   float64
7   density               1023 non-null   float64
8   pH                   1023 non-null   float64
9   sulphates             1023 non-null   float64
10  alcohol               1023 non-null   float64
11  good_quality          1023 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 96.0 KB
```

- No missing values in each feature

Validation Set

```
In [13]: df_val = pd.read_csv('/Users/chenjiqing/Public/2022_Spring_term/QBS108 Applied ML/2022 file/Homework 2/wine/val.csv')
df_val = df_val.assign(good_quality = np.where(df_val['good_quality']!=False,0,1)) # Convert target into 1, 0
df_val
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	good_quality
0	7.0	0.655	0.16	2.1	0.074	8.0	25.0	0.99606	3.37	0.55	9.7	0
1	10.4	0.440	0.42	1.5	0.145	34.0	48.0	0.99832	3.38	0.86	9.9	0
2	5.2	0.480	0.04	1.6	0.054	19.0	106.0	0.99270	3.54	0.62	12.2	1
3	9.7	0.320	0.54	2.5	0.094	28.0	83.0	0.99840	3.28	0.82	9.6	0
4	8.1	0.870	0.00	2.2	0.084	10.0	31.0	0.99656	3.25	0.50	9.8	0
...
251	8.7	0.690	0.31	3.0	0.086	23.0	81.0	1.00020	3.48	0.74	11.6	1
252	5.0	1.020	0.04	1.4	0.045	41.0	85.0	0.99380	3.75	0.48	10.5	0

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	good_quality
253	9.9	0.590	0.07	3.4	0.102	32.0	71.0	1.00015	3.31	0.71	9.8	0
254	9.1	0.500	0.30	1.9	0.065	8.0	17.0	0.99774	3.32	0.71	10.5	1
255	7.8	0.590	0.33	2.0	0.074	24.0	120.0	0.99680	3.25	0.54	9.4	0

256 rows x 12 columns

```
In [14]: df_val.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 256 entries, 0 to 255
Data columns (total 12 columns):
#   Column              Non-Null Count  Dtype
---  -
0   fixed acidity        256 non-null    float64
1   volatile acidity     256 non-null    float64
2   citric acid          256 non-null    float64
3   residual sugar       256 non-null    float64
4   chlorides            256 non-null    float64
5   free sulfur dioxide  256 non-null    float64
6   total sulfur dioxide 256 non-null    float64
7   density              256 non-null    float64
8   pH                  256 non-null    float64
9   sulphates            256 non-null    float64
10  alcohol              256 non-null    float64
11  good_quality         256 non-null    int64
dtypes: float64(11), int64(1)
memory usage: 24.1 KB
```

- No missing values in each feature

Testing Set

```
In [15]: df_test = pd.read_csv('/Users/chenjiqing/Public/2022_Spring_term/QBS108 Applied ML/2022 file/Homework 2/wine/test.csv')
df_test = df_test.assign(good_quality = np.where(df_test['good_quality']!=False,0,1)) # Convert target into 1, 0
df_test
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	good_quality
0	10.8	0.470	0.43	2.10	0.171	27.0	66.0	0.99820	3.17	0.76	10.8	1
1	8.1	0.820	0.00	4.10	0.095	5.0	14.0	0.99854	3.36	0.53	9.6	0
2	9.1	0.290	0.33	2.05	0.063	13.0	27.0	0.99516	3.26	0.84	11.7	1
3	10.2	0.645	0.36	1.80	0.053	5.0	14.0	0.99820	3.17	0.42	10.0	1
4	12.2	0.450	0.49	1.40	0.075	3.0	6.0	0.99690	3.13	0.63	10.4	0
...
315	10.1	0.270	0.54	2.30	0.065	7.0	26.0	0.99531	3.17	0.53	12.5	1
316	6.9	0.390	0.24	2.10	0.102	4.0	7.0	0.99462	3.44	0.58	11.4	0
317	9.1	0.340	0.42	1.80	0.058	9.0	18.0	0.99392	3.18	0.55	11.4	0
318	9.1	0.765	0.04	1.60	0.078	4.0	14.0	0.99800	3.29	0.54	9.7	0
319	8.2	0.320	0.42	2.30	0.098	3.0	9.0	0.99506	3.27	0.55	12.3	1

320 rows x 12 columns

```
In [16]: df_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 320 entries, 0 to 319
```

```
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          320 non-null   float64
1   volatile acidity       320 non-null   float64
2   citric acid            320 non-null   float64
3   residual sugar         320 non-null   float64
4   chlorides              320 non-null   float64
5   free sulfur dioxide    320 non-null   float64
6   total sulfur dioxide   320 non-null   float64
7   density               320 non-null   float64
8   pH                    320 non-null   float64
9   sulphates             320 non-null   float64
10  alcohol               320 non-null   float64
11  good_quality          320 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 30.1 KB
```

- No missing values in each feature

```
In [17]: # Set Features and Target for each set.
df_train_y = df_train['good_quality'].copy()
df_val_y = df_val['good_quality'].copy()
df_test_y = df_test['good_quality'].copy()
df_train = df_train.drop(['good_quality'], axis=1).copy()
df_val = df_val.drop(['good_quality'], axis=1).copy()
df_test = df_test.drop(['good_quality'], axis=1).copy()

print(df_train.shape)
print(df_train_y.shape)
print(df_val.shape)
print(df_val_y.shape)
print(df_test.shape)
print(df_test_y.shape)

(1023, 11)
(1023,)
(256, 11)
(256,)
(320, 11)
(320,)
```

```
In [18]: # Pre-processed data
from sklearn.preprocessing import StandardScaler
SS = StandardScaler()

df_train_SS = SS.fit_transform(df_train)
df_val_SS = SS.transform(df_val)
df_test_SS = SS.transform(df_test)
```

2-1. Train an k-Nearest Neighbors algorithm on the training dataset of your choice of k's

```
In [19]: from sklearn.neighbors import KNeighborsClassifier
k_cluster = []
accuracy = []
for k in range(2,11):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn = knn.fit(df_train_SS, df_train_y)
    print("k: ", k)
    print("Train Accuracy: ", knn.score(df_train_SS, df_train_y))
    print("Validation Accuracy: ", knn.score(df_val_SS, df_val_y))
    print('-'*20)
    k_cluster.append(k)
    accuracy.append(knn.score(df_val_SS, df_val_y))

k: 2
Train Accuracy: 0.8787878787878788
```

```

Validation Accuracy: 0.66796875
-----
k: 3
Train Accuracy: 0.8699902248289345
Validation Accuracy: 0.68359375
-----
k: 4
Train Accuracy: 0.8250244379276638
Validation Accuracy: 0.69140625
-----
k: 5
Train Accuracy: 0.820136852394917
Validation Accuracy: 0.71484375
-----
k: 6
Train Accuracy: 0.8005865102639296
Validation Accuracy: 0.703125
-----
k: 7
Train Accuracy: 0.7947214076246334
Validation Accuracy: 0.7265625
-----
k: 8
Train Accuracy: 0.7839687194525904
Validation Accuracy: 0.74609375
-----
k: 9
Train Accuracy: 0.7859237536656891
Validation Accuracy: 0.73828125
-----
k: 10
Train Accuracy: 0.7771260997067448
Validation Accuracy: 0.734375
-----

```

2-2. Report the classification accuracy of this model on the validation set for different values for k. Plot these accuracies against k and report the optimal value for k.

```


In [20]: plt.plot(k_cluster, accuracy)
plt.title("Accuracy on the validation set for each k")
plt.xlabel("k")
plt.ylabel("Validation Accuracy")

```

```

Out[20]: Text(0, 0.5, 'Validation Accuracy')

```



- According to the plot, the optimal value for k is 8.

2-3. Report the classification accuracy of this model on the data in test.csv using the optimal value of k that you found in 2.2.


```
In [21]: knn = KNeighborsClassifier(n_neighbors=8)
knn = knn.fit(df_train_SS, df_train_y)
print("k: ", 8)
print("Train Accuracy: ", knn.score(df_train_SS, df_train_y))
print("Testing Accuracy: ", knn.score(df_test_SS, df_test_y))
```

```
k: 8
Train Accuracy: 0.7839687194525904
Testing Accuracy: 0.71875
```

3 Decision Tree Classification

Load Data

Training Set

```
In [22]: df_train = pd.read_csv('/Users/chenjiqing/Public/2022_Spring_term/QBS108 Applied ML/2022 file/Homework 2/titanic/train.csv')
df_train
```

```
Out[22]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Cabin	Embarked	Title
0	1	1	female	24	0	0	69.3000	B35	C	Mrs
1	0	3	female	3	3	1	21.0750	NaN	S	Miss
2	1	3	female	16	0	0	7.7333	NaN	Q	Miss
3	1	3	male	32	0	0	56.4958	NaN	S	Mr
4	0	3	male	22	0	0	7.1250	NaN	S	Mr
...
493	0	2	male	52	0	0	13.5000	NaN	S	Mr
494	0	3	male	33	1	1	20.5250	NaN	S	Mr
495	1	2	male	38	0	0	13.0000	NaN	S	Mr
496	0	3	male	38	8	2	69.5500	NaN	S	Mr
497	1	1	female	56	0	1	83.1583	C50	C	Mrs

498 rows × 10 columns

```
In [23]: df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 498 entries, 0 to 497
Data columns (total 10 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Survived    498 non-null    int64
1   Pclass      498 non-null    int64
2   Sex         498 non-null    object
3   Age         498 non-null    int64
4   SibSp       498 non-null    int64
5   Parch       498 non-null    int64
6   Fare        498 non-null    float64
7   Cabin       113 non-null    object
8   Embarked    497 non-null    object
9   Title       498 non-null    object
dtypes: float64(1), int64(5), object(4)
memory usage: 39.0+ KB
```

- Cabin and Embarked contain missing values

```
In [24]: print("Sex:")
```

```
print(df_train.Sex.value_counts(), "\n")
print("Cabin:")
print(df_train.Cabin.value_counts(), "\n")
print("Embarked:")
print(df_train.Embarked.value_counts(), "\n")
print("Title:")
print(df_train.Title.value_counts(), "\n")
```

Sex:
male 333
female 165
Name: Sex, dtype: int64

Cabin:
C23 C25 C27 4
F101 3
G6 3
D35 2
C22 C26 2
..
B4 1
C78 1
D30 1
C87 1
B37 1
Name: Cabin, Length: 91, dtype: int64

Embarked:
S 373
C 85
Q 39
Name: Embarked, dtype: int64

Title:
Mr 295
Miss 102
Mrs 61
Master 27
Rare 13
Name: Title, dtype: int64

Validation Set

```
In [25]: df_val = pd.read_csv('/Users/chenjiqing/Public/2022_Spring_term/QBS108 Applied ML/2022 file/Homework 2/titanic/val.csv')
df_val
```

Out[25]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Cabin	Embarked	Title
0	1	3	female	24	1	0	15.8500	NaN	S	Mrs
1	1	2	female	24	0	0	13.0000	F33	S	Miss
2	0	3	male	40	1	1	15.5000	NaN	Q	Mr
3	1	3	male	22	2	0	23.2500	NaN	Q	Mr
4	1	2	female	3	1	2	41.5792	NaN	C	Miss
...
120	1	2	female	31	1	1	26.2500	NaN	S	Mrs
121	1	3	male	19	0	0	7.7500	NaN	Q	Mr
122	1	1	female	43	0	1	211.3375	B3	S	Mrs
123	0	2	male	52	0	0	13.0000	NaN	S	Mr
124	1	1	female	35	1	0	53.1000	C123	S	Mrs

125 rows x 10 columns

In [26]: `df_val.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 125 entries, 0 to 124
Data columns (total 10 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Survived    125 non-null    int64
1   Pclass      125 non-null    int64
2   Sex         125 non-null    object
3   Age         125 non-null    int64
4   SibSp       125 non-null    int64
5   Parch       125 non-null    int64
6   Fare        125 non-null    float64
7   Cabin       26 non-null     object
8   Embarked    125 non-null    object
9   Title       125 non-null    object
dtypes: float64(1), int64(5), object(4)
memory usage: 9.9+ KB
```

- Cabin contains missing values

In [27]:

```
print("Sex:")
print(df_val.Sex.value_counts(), "\n")
print("Cabin:")
print(df_val.Cabin.value_counts(), "\n")
print("Embarked:")
print(df_val.Embarked.value_counts(), "\n")
print("Title:")
print(df_val.Title.value_counts(), "\n")
```

```
Sex:
male      77
female    48
Name: Sex, dtype: int64
```

```
Cabin:
C65      2
C83      2
D46      1
D37      1
B101     1
B71      1
B51 B53 B55 1
E121     1
F33      1
C2       1
F38      1
E33      1
B3       1
B18      1
C124     1
D9       1
D26      1
C125     1
C123     1
C68      1
A19      1
E58      1
D49      1
A24      1
Name: Cabin, dtype: int64
```

```
Embarked:
S      84
C      25
Q      16
Name: Embarked, dtype: int64
```

```
Title:
```

```
Mr      73
Miss    26
Mrs     22
Master   4
Name: Title, dtype: int64
```

Testing Set

```
In [28]: df_test = pd.read_csv('/Users/chenjiqing/Public/2022_Spring_term/QBS108 Applied ML/2022 file/Homework 2/titanic/test.csv')
df_test
```

```
Out[28]:
```

	Pclass	Sex	Age	SibSp	Parch	Fare	Cabin	Embarked	Title
0	3	male	22	1	1	15.2458	NaN	C	Master
1	2	male	31	0	0	10.5000	NaN	S	Mr
2	3	male	20	0	0	7.9250	NaN	S	Mr
3	2	female	6	0	1	33.0000	NaN	S	Miss
4	3	female	14	1	0	11.2417	NaN	C	Miss
...
263	3	male	27	0	0	8.6625	NaN	S	Mr
264	1	male	27	0	0	0.0000	NaN	S	Mr
265	1	male	65	0	0	26.5500	E38	S	Mr
266	3	male	17	0	0	8.6625	NaN	S	Mr
267	1	male	28	0	0	26.5500	C52	S	Mr

268 rows x 9 columns

```
In [29]: df_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 268 entries, 0 to 267
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Pclass      268 non-null    int64
1   Sex         268 non-null    object
2   Age         268 non-null    int64
3   SibSp       268 non-null    int64
4   Parch       268 non-null    int64
5   Fare        268 non-null    float64
6   Cabin       65 non-null     object
7   Embarked    267 non-null    object
8   Title       268 non-null    object
dtypes: float64(1), int64(4), object(4)
memory usage: 19.0+ KB
```

- Cabin and Embarked contain missing values

```
In [30]: print("Sex:")
print(df_test.Sex.value_counts(), "\n")
print("Cabin:")
print(df_test.Cabin.value_counts(), "\n")
print("Embarked:")
print(df_test.Embarked.value_counts(), "\n")
print("Title:")
print(df_test.Title.value_counts(), "\n")
```

```
Sex:
male      167
female    101
```

Name: Sex, dtype: int64

Cabin:

E25	2
B58 B60	2
D	2
B57 B59 B63 B66	2
C126	2
B96 B98	2
C52	2
C92	1
B49	1
E34	1
C82	1
C148	1
C110	1
D17	1
C104	1
C123	1
D11	1
C86	1
A31	1
A7	1
E68	1
B69	1
D21	1
D19	1
D10 D12	1
B28	1
E77	1
A6	1
E49	1
E36	1
D20	1
B39	1
C78	1
C45	1
B30	1
F33	1
E63	1
G6	1
B86	1
C22 C26	1
F2	1
E38	1
A20	1
B38	1
D7	1
D28	1
D47	1
D36	1
D15	1
C106	1
C95	1
C124	1
A5	1
D48	1
B20	1
D50	1
B78	1
C93	1

Name: Cabin, dtype: int64

Embarked:

S	187
C	58
Q	22

Name: Embarked, dtype: int64

Title:

Mr	149
Miss	57

```

Mrs      43
Rare     10
Master   9
Name: Title, dtype: int64

```

3-1. Preprocess the data

- Set Features and Target for training and validation set.

```

In [31]: df_train_y = df_train['Survived'].copy()
df_val_y = df_val['Survived'].copy()
df_train = df_train.drop(['Survived'], axis=1).copy()
df_val = df_val.drop(['Survived'], axis=1).copy()

print(df_train.shape)
print(df_train_y.shape)
print(df_val.shape)
print(df_val_y.shape)

(498, 9)
(498,)
(125, 9)
(125,)

```

- For each set, more than 75% values in Cabin are missing values so I will remove this feature for the following analysis.

```

In [32]: df_train2 = df_train.drop('Cabin',axis=1).copy()
df_val2 = df_val.drop('Cabin',axis=1).copy()
df_test2 = df_test.drop('Cabin',axis=1).copy()

```

- Because Decision Tree only can handle numbers so I need to convert the strings to numeric. However, in vlidation set, the 'Title' feature doesn't have 'Rare' label. As a result, I need to combine training, testing, and validation set together first, then, convert the strings to numeric.

```

In [33]: # Combine 3 dataset
df_train2['set'] = "Train"
df_val2['set'] = "Val"
df_test2['set'] = "Test"

df = pd.concat([df_train2,df_val2,df_test2])
df

```

```

Out[33]:

```

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	Title	set
0	1	female	24	0	0	69.3000	C	Mrs	Train
1	3	female	3	3	1	21.0750	S	Miss	Train
2	3	female	16	0	0	7.7333	Q	Miss	Train
3	3	male	32	0	0	56.4958	S	Mr	Train
4	3	male	22	0	0	7.1250	S	Mr	Train
...
263	3	male	27	0	0	8.6625	S	Mr	Test
264	1	male	27	0	0	0.0000	S	Mr	Test
265	1	male	65	0	0	26.5500	S	Mr	Test
266	3	male	17	0	0	8.6625	S	Mr	Test
267	1	male	28	0	0	26.5500	S	Mr	Test

891 rows × 9 columns

```
In [34]: print("Sex:")
print(df.Sex.value_counts(), "\n")
print("Embarked:")
print(df.Embarked.value_counts(), "\n")
print("Title:")
print(df.Title.value_counts(), "\n")
```

```
Sex:
male      577
female    314
Name: Sex, dtype: int64
```

```
Embarked:
S      644
C     168
Q       77
Name: Embarked, dtype: int64
```

```
Title:
Mr      517
Miss    185
Mrs     126
Master   40
Rare     23
Name: Title, dtype: int64
```

```
In [35]: # convert the strings to numeric
# Since Embarked has missing values, I will turn missing into new category.
df = df.assign(Sex = np.where(df['Sex']=='male',0,1),
               Embarked = np.where(df['Embarked']=='C',0,np.where(df['Embarked']=='S',1,
               np.where(df['Embarked']=='Q',2,3))),
               Title = np.where(df['Title']=='Mr',0,np.where(df['Title']=='Miss',1,
               np.where(df['Title']=='Mrs',2,
               np.where(df['Title']=='Master',3,4))))).copy()

df
```

```
Out[35]:
```

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	Title	set
0	1	1	24	0	0	69.3000	0	2	Train
1	3	1	3	3	1	21.0750	1	1	Train
2	3	1	16	0	0	7.7333	2	1	Train
3	3	0	32	0	0	56.4958	1	0	Train
4	3	0	22	0	0	7.1250	1	0	Train
...
263	3	0	27	0	0	8.6625	1	0	Test
264	1	0	27	0	0	0.0000	1	0	Test
265	1	0	65	0	0	26.5500	1	0	Test
266	3	0	17	0	0	8.6625	1	0	Test
267	1	0	28	0	0	26.5500	1	0	Test

891 rows × 9 columns

```
In [36]: # Check
print("Sex:")
print(df.Sex.value_counts(), "\n")
print("Embarked:")
print(df.Embarked.value_counts(), "\n")
```

```
print("Title:")
print(df.Title.value_counts(), "\n")
```

```
Sex:
0    577
1    314
Name: Sex, dtype: int64
```

```
Embarked:
1    644
0    168
2     77
3      2
Name: Embarked, dtype: int64
```

```
Title:
0    517
1    185
2    126
3     40
4     23
Name: Title, dtype: int64
```

```
In [37]: # Split it back to 3 dataset
df_train3 = df[df['set'] == 'Train'].copy()
df_train3 = df_train3.drop('set', axis=1)
print(df_train3)

df_val3 = df[df['set'] == 'Val'].copy()
df_val3 = df_val3.drop('set', axis=1)
print(df_val3)

df_test3 = df[df['set'] == 'Test'].copy()
df_test3 = df_test3.drop('set', axis=1)
print(df_test3)
```

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	Title
0	1	1	24	0	0	69.3000	0	2
1	3	1	3	3	1	21.0750	1	1
2	3	1	16	0	0	7.7333	2	1
3	3	0	32	0	0	56.4958	1	0
4	3	0	22	0	0	7.1250	1	0
..
493	2	0	52	0	0	13.5000	1	0
494	3	0	33	1	1	20.5250	1	0
495	2	0	38	0	0	13.0000	1	0
496	3	0	38	8	2	69.5500	1	0
497	1	1	56	0	1	83.1583	0	2

[498 rows x 8 columns]

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	Title
0	3	1	24	1	0	15.8500	1	2
1	2	1	24	0	0	13.0000	1	1
2	3	0	40	1	1	15.5000	2	0
3	3	0	22	2	0	23.2500	2	0
4	2	1	3	1	2	41.5792	0	1
..
120	2	1	31	1	1	26.2500	1	2
121	3	0	19	0	0	7.7500	2	0
122	1	1	43	0	1	211.3375	1	2
123	2	0	52	0	0	13.0000	1	0
124	1	1	35	1	0	53.1000	1	2

[125 rows x 8 columns]

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	Title
0	3	0	22	1	1	15.2458	0	3
1	2	0	31	0	0	10.5000	1	0
2	3	0	20	0	0	7.9250	1	0
3	2	1	6	0	1	33.0000	1	1
4	3	1	14	1	0	11.2417	0	1
..

263	3	0	27	0	0	8.6625	1	0
264	1	0	27	0	0	0.0000	1	0
265	1	0	65	0	0	26.5500	1	0
266	3	0	17	0	0	8.6625	1	0
267	1	0	28	0	0	26.5500	1	0

[268 rows x 8 columns]

3-2. Decision Trees Classification

3-2-a. Initialize a binary decision tree model for the training data.

```
In [38]: from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
dt = DecisionTreeClassifier(random_state=0)
```

```
In [39]: dt2 = dt.fit(df_train3, df_train_y)
tree.plot_tree(dt2)
```

```
Out[39]: [Text(173.75830792682927, 211.04470588235293, 'X[7] <= 0.5\ngini = 0.471\nsamples = 498\nvalue = [309, 189]'),
Text(83.94100609756097, 198.25411764705882, 'X[0] <= 1.5\ngini = 0.277\nsamples = 295\nvalue = [246, 49]'),
Text(19.053658536585367, 185.4635294117647, 'X[5] <= 15.644\ngini = 0.46\nsamples = 53\nvalue = [34, 19]'),
Text(15.42439024390244, 172.6729411764706, 'gini = 0.0\nsamples = 4\nvalue = [4, 0]'),
Text(22.682926829268293, 172.6729411764706, 'X[5] <= 27.135\ngini = 0.475\nsamples = 49\nvalue = [30, 19]'),
Text(10.887804878048781, 159.88235294117646, 'X[2] <= 53.5\ngini = 0.408\nsamples = 7\nvalue = [2, 5]'),
Text(7.258536585365854, 147.09176470588235, 'gini = 0.0\nsamples = 5\nvalue = [0, 5]'),
Text(14.517073170731708, 147.09176470588235, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(34.4780487804878, 159.88235294117646, 'X[3] <= 0.5\ngini = 0.444\nsamples = 42\nvalue = [28, 14]'),
Text(21.775609756097563, 147.09176470588235, 'X[4] <= 1.5\ngini = 0.34\nsamples = 23\nvalue = [18, 5]'),
Text(18.146341463414636, 134.30117647058825, 'X[5] <= 32.0\ngini = 0.298\nsamples = 22\nvalue = [18, 4]'),
Text(14.517073170731708, 121.51058823529411, 'X[5] <= 29.85\ngini = 0.48\nsamples = 10\nvalue = [6, 4]'),
Text(7.258536585365854, 108.72, 'X[5] <= 28.725\ngini = 0.32\nsamples = 5\nvalue = [4, 1]'),
Text(3.629268292682927, 95.92941176470588, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(10.887804878048781, 95.92941176470588, 'X[4] <= 0.5\ngini = 0.444\nsamples = 3\nvalue = [2, 1]'),
Text(7.258536585365854, 83.13882352941175, 'X[2] <= 48.0\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(3.629268292682927, 70.34823529411764, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(10.887804878048781, 70.34823529411764, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(14.517073170731708, 83.13882352941175, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(21.775609756097563, 108.72, 'X[2] <= 28.0\ngini = 0.48\nsamples = 5\nvalue = [2, 3]'),
Text(18.146341463414636, 95.92941176470588, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(25.40487804878049, 95.92941176470588, 'X[5] <= 30.25\ngini = 0.5\nsamples = 4\nvalue = [2, 2]'),
Text(21.775609756097563, 83.13882352941175, 'X[2] <= 33.0\ngini = 0.444\nsamples = 3\nvalue = [1, 2]'),
Text(18.146341463414636, 70.34823529411764, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(25.40487804878049, 70.34823529411764, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
Text(29.034146341463416, 83.13882352941175, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(21.775609756097563, 121.51058823529411, 'gini = 0.0\nsamples = 12\nvalue = [12, 0]'),
Text(25.40487804878049, 134.30117647058825, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(47.18048780487805, 147.09176470588235, 'X[5] <= 100.981\ngini = 0.499\nsamples = 19\nvalue = [10, 9]'),
Text(43.551219512195125, 134.30117647058825, 'X[5] <= 85.638\ngini = 0.492\nsamples = 16\nvalue = [7, 9]'),
Text(39.921951219512195, 121.51058823529411, 'X[5] <= 79.425\ngini = 0.497\nsamples = 13\nvalue = [7, 6]'),
Text(36.29268292682927, 108.72, 'X[6] <= 0.5\ngini = 0.496\nsamples = 11\nvalue = [5, 6]'),
Text(32.66341463414634, 95.92941176470588, 'gini = 0.0\nsamples = 3\nvalue = [0, 3]'),
Text(39.921951219512195, 95.92941176470588, 'X[2] <= 39.5\ngini = 0.469\nsamples = 8\nvalue = [5, 3]'),
Text(36.29268292682927, 83.13882352941175, 'X[5] <= 67.925\ngini = 0.48\nsamples = 5\nvalue = [2, 3]'),
Text(32.66341463414634, 70.34823529411764, 'X[2] <= 23.0\ngini = 0.375\nsamples = 4\nvalue = [1, 3]'),
Text(29.034146341463416, 57.557647058823534, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(36.29268292682927, 57.557647058823534, 'gini = 0.0\nsamples = 3\nvalue = [0, 3]'),
Text(39.921951219512195, 70.34823529411764, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(43.551219512195125, 83.13882352941175, 'gini = 0.0\nsamples = 3\nvalue = [3, 0]'),
Text(43.551219512195125, 108.72, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(47.18048780487805, 121.51058823529411, 'gini = 0.0\nsamples = 3\nvalue = [0, 3]'),
Text(50.80975609756098, 134.30117647058825, 'gini = 0.0\nsamples = 3\nvalue = [3, 0]'),
Text(148.8283536585366, 185.4635294117647, 'X[5] <= 51.698\ngini = 0.217\nsamples = 242\nvalue = [212, 30]'),
Text(106.21280487804879, 172.6729411764706, 'X[6] <= 0.5\ngini = 0.205\nsamples = 233\nvalue = [206, 27]'),
Text(68.9560975609756, 159.88235294117646, 'X[2] <= 29.5\ngini = 0.337\nsamples = 28\nvalue = [22, 6]'),
Text(65.32682926829268, 147.09176470588235, 'X[2] <= 28.5\ngini = 0.469\nsamples = 16\nvalue = [10, 6]'),
Text(61.69756097560976, 134.30117647058825, 'X[5] <= 10.546\ngini = 0.408\nsamples = 14\nvalue = [10, 4]'),
Text(54.4390243902439, 121.51058823529411, 'X[2] <= 22.5\ngini = 0.198\nsamples = 9\nvalue = [8, 1]'),
Text(50.80975609756098, 108.72, 'X[2] <= 21.0\ngini = 0.375\nsamples = 4\nvalue = [3, 1]'),
```


Text(47.18048780487805, 95.92941176470588, 'gini = 0.0\nsamples = 3\nvalue = [3, 0]'),
Text(54.4390243902439, 95.92941176470588, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(58.06829268292683, 108.72, 'gini = 0.0\nsamples = 5\nvalue = [5, 0]'),
Text(68.9560975609756, 121.51058823529411, 'X[2] <= 23.5\ngini = 0.48\nsamples = 5\nvalue = [2, 3]'),
Text(65.32682926829268, 108.72, 'X[4] <= 0.5\ngini = 0.444\nsamples = 3\nvalue = [2, 1]'),
Text(61.69756097560976, 95.92941176470588, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(68.9560975609756, 95.92941176470588, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(72.58536585365854, 108.72, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
Text(68.9560975609756, 134.30117647058825, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
Text(72.58536585365854, 147.09176470588235, 'gini = 0.0\nsamples = 12\nvalue = [12, 0]'),
Text(143.46951219512195, 159.88235294117646, 'X[5] <= 13.25\ngini = 0.184\nsamples = 205\nvalue = [184, 21]'),
Text(117.27073170731708, 147.09176470588235, 'X[2] <= 30.5\ngini = 0.22\nsamples = 159\nvalue = [139, 20]'),
Text(96.62926829268292, 134.30117647058825, 'X[5] <= 8.206\ngini = 0.173\nsamples = 94\nvalue = [85, 9]'),
Text(93.0, 121.51058823529411, 'X[5] <= 8.081\ngini = 0.242\nsamples = 64\nvalue = [55, 9]'),
Text(89.37073170731708, 108.72, 'X[3] <= 0.5\ngini = 0.222\nsamples = 63\nvalue = [55, 8]'),
Text(76.21463414634147, 95.92941176470588, 'X[5] <= 7.825\ngini = 0.206\nsamples = 60\nvalue = [53, 7]'),
Text(60.79024390243902, 83.13882352941175, 'X[2] <= 23.0\ngini = 0.278\nsamples = 30\nvalue = [25, 5]'),
Text(51.71707317073171, 70.34823529411764, 'X[5] <= 7.785\ngini = 0.105\nsamples = 18\nvalue = [17, 1]'),
Text(48.08780487804878, 57.557647058823534, 'gini = 0.0\nsamples = 16\nvalue = [16, 0]'),
Text(55.34634146341464, 57.557647058823534, 'X[2] <= 21.5\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(51.71707317073171, 44.767058823529396, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(58.97560975609756, 44.767058823529396, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(69.86341463414634, 70.34823529411764, 'X[5] <= 3.525\ngini = 0.444\nsamples = 12\nvalue = [8, 4]'),
Text(66.23414634146341, 57.557647058823534, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(73.49268292682927, 57.557647058823534, 'X[2] <= 26.5\ngini = 0.397\nsamples = 11\nvalue = [8, 3]'),
Text(66.23414634146341, 44.767058823529396, 'X[5] <= 7.396\ngini = 0.278\nsamples = 6\nvalue = [5, 1]'),
Text(62.60487804878049, 31.976470588235287, 'X[2] <= 24.5\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(58.97560975609756, 19.185882352941178, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(66.23414634146341, 19.185882352941178, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(69.86341463414634, 31.976470588235287, 'gini = 0.0\nsamples = 4\nvalue = [4, 0]'),
Text(80.75121951219512, 44.767058823529396, 'X[2] <= 27.5\ngini = 0.48\nsamples = 5\nvalue = [3, 2]'),
Text(77.1219512195122, 31.976470588235287, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(84.38048780487806, 31.976470588235287, 'X[6] <= 1.5\ngini = 0.375\nsamples = 4\nvalue = [3, 1]'),
Text(80.75121951219512, 19.185882352941178, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(88.00975609756098, 19.185882352941178, 'X[2] <= 29.5\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(84.38048780487806, 6.39529411764704, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(91.6390243902439, 6.39529411764704, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(91.6390243902439, 83.13882352941175, 'X[2] <= 19.5\ngini = 0.124\nsamples = 30\nvalue = [28, 2]'),
Text(88.00975609756098, 70.34823529411764, 'X[5] <= 7.973\ngini = 0.444\nsamples = 6\nvalue = [4, 2]'),
Text(84.38048780487806, 57.557647058823534, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(91.6390243902439, 57.557647058823534, 'X[2] <= 16.5\ngini = 0.48\nsamples = 5\nvalue = [3, 2]'),
Text(88.00975609756098, 44.767058823529396, 'gini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(95.26829268292683, 44.767058823529396, 'X[2] <= 18.0\ngini = 0.444\nsamples = 3\nvalue = [2, 1]'),
Text(91.6390243902439, 31.976470588235287, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(98.89756097560976, 31.976470588235287, 'gini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(95.26829268292683, 70.34823529411764, 'gini = 0.0\nsamples = 24\nvalue = [24, 0]'),
Text(102.52682926829269, 95.92941176470588, 'X[2] <= 21.0\ngini = 0.444\nsamples = 3\nvalue = [2, 1]'),
Text(98.89756097560976, 83.13882352941175, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(106.15609756097561, 83.13882352941175, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(96.62926829268292, 108.72, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(100.25853658536586, 121.51058823529411, 'gini = 0.0\nsamples = 30\nvalue = [30, 0]'),
Text(137.9121951219512, 134.30117647058825, 'X[5] <= 7.763\ngini = 0.281\nsamples = 65\nvalue = [54, 11]'),
Text(134.2829268292683, 121.51058823529411, 'gini = 0.0\nsamples = 20\nvalue = [20, 0]'),
Text(141.54146341463414, 121.51058823529411, 'X[2] <= 32.5\ngini = 0.369\nsamples = 45\nvalue = [34, 11]'),
Text(124.30243902439025, 108.72, 'X[5] <= 7.91\ngini = 0.5\nsamples = 8\nvalue = [4, 4]'),
Text(117.0439024390244, 95.92941176470588, 'X[5] <= 7.875\ngini = 0.375\nsamples = 4\nvalue = [3, 1]'),
Text(113.41463414634147, 83.13882352941175, 'X[2] <= 31.5\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(109.78536585365855, 70.34823529411764, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(117.0439024390244, 70.34823529411764, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(120.67317073170732, 83.13882352941175, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(131.5609756097561, 95.92941176470588, 'X[5] <= 8.206\ngini = 0.375\nsamples = 4\nvalue = [1, 3]'),
Text(127.93170731707318, 83.13882352941175, 'gini = 0.0\nsamples = 3\nvalue = [0, 3]'),
Text(135.19024390243902, 83.13882352941175, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(158.78048780487805, 108.72, 'X[5] <= 10.0\ngini = 0.307\nsamples = 37\nvalue = [30, 7]'),
Text(146.07804878048782, 95.92941176470588, 'X[2] <= 38.5\ngini = 0.191\nsamples = 28\nvalue = [25, 3]'),
Text(142.4487804878049, 83.13882352941175, 'gini = 0.0\nsamples = 15\nvalue = [15, 0]'),
Text(149.70731707317074, 83.13882352941175, 'X[5] <= 7.988\ngini = 0.355\nsamples = 13\nvalue = [10, 3]'),
Text(146.07804878048782, 70.34823529411764, 'X[2] <= 46.0\ngini = 0.49\nsamples = 7\nvalue = [4, 3]'),
Text(142.4487804878049, 57.557647058823534, 'X[5] <= 7.835\ngini = 0.48\nsamples = 5\nvalue = [2, 3]'),
Text(138.81951219512194, 44.767058823529396, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(146.07804878048782, 44.767058823529396, 'X[5] <= 7.91\ngini = 0.5\nsamples = 4\nvalue = [2, 2]'),
Text(142.4487804878049, 31.976470588235287, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),

Text(149.70731707317074, 31.976470588235287, 'X[2] <= 41.5\ngini = 0.444\nsamples = 3\nvalue = [1, 2]'),
Text(146.07804878048782, 19.185882352941178, 'gini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(153.33658536585367, 19.185882352941178, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(149.70731707317074, 57.557647058823534, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(153.33658536585367, 70.34823529411764, 'gini = 0.0\nsamples = 6\nvalue = [6, 0]'),
Text(171.4829268292683, 95.92941176470588, 'X[2] <= 45.0\ngini = 0.494\nsamples = 9\nvalue = [5, 4]'),
Text(164.22439024390243, 83.13882352941175, 'X[2] <= 40.5\ngini = 0.48\nsamples = 5\nvalue = [2, 3]'),
Text(160.5951219512195, 70.34823529411764, 'X[2] <= 38.5\ngini = 0.5\nsamples = 4\nvalue = [2, 2]'),
Text(156.9658536585366, 57.557647058823534, 'X[2] <= 36.0\ngini = 0.444\nsamples = 3\nvalue = [1, 2]'),
Text(153.33658536585367, 44.767058823529396, 'gini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(160.5951219512195, 44.767058823529396, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(164.22439024390243, 57.557647058823534, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(167.85365853658539, 70.34823529411764, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(178.74146341463415, 83.13882352941175, 'X[5] <= 11.75\ngini = 0.375\nsamples = 4\nvalue = [3, 1]'),
Text(175.11219512195123, 70.34823529411764, 'X[2] <= 66.0\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(171.4829268292683, 57.557647058823534, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(178.74146341463415, 57.557647058823534, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(182.37073170731708, 70.34823529411764, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(169.66829268292685, 147.09176470588235, 'X[5] <= 25.075\ngini = 0.043\nsamples = 46\nvalue = [45, 1]'),
Text(166.0390243902439, 134.30117647058825, 'gini = 0.0\nsamples = 29\nvalue = [29, 0]'),
Text(173.29756097560977, 134.30117647058825, 'X[2] <= 33.0\ngini = 0.111\nsamples = 17\nvalue = [16, 1]'),
Text(169.66829268292685, 121.51058823529411, 'X[2] <= 31.5\ngini = 0.245\nsamples = 7\nvalue = [6, 1]'),
Text(166.0390243902439, 108.72, 'gini = 0.0\nsamples = 6\nvalue = [6, 0]'),
Text(173.29756097560977, 108.72, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(176.9268292682927, 121.51058823529411, 'gini = 0.0\nsamples = 10\nvalue = [10, 0]'),
Text(191.44390243902438, 172.6729411764706, 'X[5] <= 63.023\ngini = 0.444\nsamples = 9\nvalue = [6, 3]'),
Text(187.81463414634146, 159.88235294117646, 'X[2] <= 30.0\ngini = 0.375\nsamples = 4\nvalue = [1, 3]'),
Text(184.18536585365854, 147.09176470588235, 'X[2] <= 27.0\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(180.55609756097562, 134.30117647058825, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(187.81463414634146, 134.30117647058825, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(191.44390243902438, 147.09176470588235, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
Text(195.07317073170734, 159.88235294117646, 'gini = 0.0\nsamples = 5\nvalue = [5, 0]'),
Text(263.57560975609755, 198.25411764705882, 'X[0] <= 2.5\ngini = 0.428\nsamples = 203\nvalue = [63, 140]'),
Text(228.6439024390244, 185.4635294117647, 'X[7] <= 3.5\ngini = 0.215\nsamples = 106\nvalue = [13, 93]'),
Text(215.0341463414634, 172.6729411764706, 'X[5] <= 149.035\ngini = 0.082\nsamples = 93\nvalue = [4, 89]'),
Text(205.9609756097561, 159.88235294117646, 'X[2] <= 49.5\ngini = 0.048\nsamples = 82\nvalue = [2, 80]'),
Text(198.70243902439026, 147.09176470588235, 'X[5] <= 22.0\ngini = 0.026\nsamples = 75\nvalue = [1, 74]'),
Text(195.07317073170734, 134.30117647058825, 'X[5] <= 20.25\ngini = 0.087\nsamples = 22\nvalue = [1, 21]'),
Text(191.44390243902438, 121.51058823529411, 'gini = 0.0\nsamples = 20\nvalue = [0, 20]'),
Text(198.70243902439026, 121.51058823529411, 'X[2] <= 28.5\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(195.07317073170734, 108.72, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(202.33170731707318, 108.72, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(202.33170731707318, 134.30117647058825, 'gini = 0.0\nsamples = 53\nvalue = [0, 53]'),
Text(213.21951219512195, 147.09176470588235, 'X[2] <= 51.5\ngini = 0.245\nsamples = 7\nvalue = [1, 6]'),
Text(209.59024390243903, 134.30117647058825, 'X[5] <= 19.606\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(205.9609756097561, 121.51058823529411, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(213.21951219512195, 121.51058823529411, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(216.84878048780487, 134.30117647058825, 'gini = 0.0\nsamples = 5\nvalue = [0, 5]'),
Text(224.10731707317075, 159.88235294117646, 'X[5] <= 152.506\ngini = 0.298\nsamples = 11\nvalue = [2, 9]'),
Text(220.47804878048782, 147.09176470588235, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(227.73658536585367, 147.09176470588235, 'gini = 0.0\nsamples = 9\nvalue = [0, 9]'),
Text(242.25365853658536, 172.6729411764706, 'X[1] <= 0.5\ngini = 0.426\nsamples = 13\nvalue = [9, 4]'),
Text(238.62439024390244, 159.88235294117646, 'X[6] <= 0.5\ngini = 0.298\nsamples = 11\nvalue = [9, 2]'),
Text(234.99512195121952, 147.09176470588235, 'X[5] <= 37.55\ngini = 0.444\nsamples = 3\nvalue = [1, 2]'),
Text(231.3658536585366, 134.30117647058825, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
Text(238.62439024390244, 134.30117647058825, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(242.25365853658536, 147.09176470588235, 'gini = 0.0\nsamples = 8\nvalue = [8, 0]'),
Text(245.8829268292683, 159.88235294117646, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
Text(298.5073170731707, 185.4635294117647, 'X[5] <= 22.904\ngini = 0.5\nsamples = 97\nvalue = [50, 47]'),
Text(269.47317073170734, 172.6729411764706, 'X[7] <= 2.5\ngini = 0.469\nsamples = 72\nvalue = [27, 45]'),
Text(265.8439024390244, 159.88235294117646, 'X[5] <= 7.744\ngini = 0.483\nsamples = 66\nvalue = [27, 39]'),
Text(249.51219512195124, 147.09176470588235, 'X[5] <= 6.987\ngini = 0.18\nsamples = 10\nvalue = [1, 9]'),
Text(245.8829268292683, 134.30117647058825, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(253.14146341463416, 134.30117647058825, 'gini = 0.0\nsamples = 9\nvalue = [0, 9]'),
Text(282.1756097560976, 147.09176470588235, 'X[2] <= 26.5\ngini = 0.497\nsamples = 56\nvalue = [26, 30]'),
Text(260.40000000000003, 134.30117647058825, 'X[4] <= 0.5\ngini = 0.48\nsamples = 30\nvalue = [18, 12]'),
Text(247.69756097560978, 121.51058823529411, 'X[6] <= 1.5\ngini = 0.305\nsamples = 16\nvalue = [13, 3]'),
Text(240.4390243902439, 108.72, 'X[2] <= 25.5\ngini = 0.142\nsamples = 13\nvalue = [12, 1]'),
Text(236.80975609756098, 95.92941176470588, 'gini = 0.0\nsamples = 11\nvalue = [11, 0]'),
Text(244.06829268292682, 95.92941176470588, 'X[7] <= 1.5\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(240.4390243902439, 83.13882352941175, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(247.69756097560978, 83.13882352941175, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),

```

Text(254.95609756097562, 108.72, 'X[5] <= 7.89\ngini = 0.444\nsamples = 3\nvalue = [1, 2]'),
Text(251.3268292682927, 95.92941176470588, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(258.5853658536586, 95.92941176470588, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
Text(273.10243902439026, 121.51058823529411, 'X[7] <= 1.5\ngini = 0.459\nsamples = 14\nvalue = [5, 9]'),
Text(269.47317073170734, 108.72, 'X[2] <= 17.5\ngini = 0.496\nsamples = 11\nvalue = [5, 6]'),
Text(265.8439024390244, 95.92941176470588, 'X[3] <= 2.5\ngini = 0.444\nsamples = 9\nvalue = [3, 6]'),
Text(258.5853658536586, 83.13882352941175, 'X[5] <= 11.375\ngini = 0.278\nsamples = 6\nvalue = [1, 5]'),
Text(254.95609756097562, 70.34823529411764, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(262.2146341463415, 70.34823529411764, 'gini = 0.0\nsamples = 5\nvalue = [0, 5]'),
Text(273.10243902439026, 83.13882352941175, 'X[5] <= 14.5\ngini = 0.444\nsamples = 3\nvalue = [2, 1]'),
Text(269.47317073170734, 70.34823529411764, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(276.7317073170732, 70.34823529411764, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(273.10243902439026, 95.92941176470588, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(276.7317073170732, 108.72, 'gini = 0.0\nsamples = 3\nvalue = [0, 3]'),
Text(303.9512195121951, 134.30117647058825, 'X[6] <= 1.5\ngini = 0.426\nsamples = 26\nvalue = [8, 18]'),
Text(294.8780487804878, 121.51058823529411, 'X[2] <= 36.5\ngini = 0.484\nsamples = 17\nvalue = [7, 10]'),
Text(287.61951219512196, 108.72, 'X[4] <= 3.0\ngini = 0.397\nsamples = 11\nvalue = [3, 8]'),
Text(283.99024390243903, 95.92941176470588, 'X[3] <= 0.5\ngini = 0.32\nsamples = 10\nvalue = [2, 8]'),
Text(280.3609756097561, 83.13882352941175, 'gini = 0.0\nsamples = 4\nvalue = [0, 4]'),
Text(287.61951219512196, 83.13882352941175, 'X[2] <= 32.0\ngini = 0.444\nsamples = 6\nvalue = [2, 4]'),
Text(283.99024390243903, 70.34823529411764, 'X[5] <= 19.262\ngini = 0.444\nsamples = 3\nvalue = [2, 1]'),
Text(280.3609756097561, 57.557647058823534, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(287.61951219512196, 57.557647058823534, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(291.2487804878049, 70.34823529411764, 'gini = 0.0\nsamples = 3\nvalue = [0, 3]'),
Text(291.2487804878049, 95.92941176470588, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(302.13658536585365, 108.72, 'X[5] <= 15.673\ngini = 0.444\nsamples = 6\nvalue = [4, 2]'),
Text(298.5073170731707, 95.92941176470588, 'gini = 0.0\nsamples = 4\nvalue = [4, 0]'),
Text(305.76585365853657, 95.92941176470588, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
Text(313.0243902439025, 121.51058823529411, 'X[2] <= 41.0\ngini = 0.198\nsamples = 9\nvalue = [1, 8]'),
Text(309.3951219512195, 108.72, 'gini = 0.0\nsamples = 7\nvalue = [0, 7]'),
Text(316.6536585365854, 108.72, 'X[3] <= 0.5\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(313.0243902439025, 95.92941176470588, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(320.2829268292683, 95.92941176470588, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(273.10243902439026, 159.88235294117646, 'gini = 0.0\nsamples = 6\nvalue = [0, 6]'),
Text(327.54146341463417, 172.6729411764706, 'X[2] <= 6.0\ngini = 0.147\nsamples = 25\nvalue = [23, 2]'),
Text(323.91219512195124, 159.88235294117646, 'X[2] <= 4.5\ngini = 0.346\nsamples = 9\nvalue = [7, 2]'),
Text(320.2829268292683, 147.09176470588235, 'X[2] <= 2.5\ngini = 0.219\nsamples = 8\nvalue = [7, 1]'),
Text(316.6536585365854, 134.30117647058825, 'gini = 0.0\nsamples = 5\nvalue = [5, 0]'),
Text(323.91219512195124, 134.30117647058825, 'X[2] <= 3.5\ngini = 0.444\nsamples = 3\nvalue = [2, 1]'),
Text(320.2829268292683, 121.51058823529411, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(327.54146341463417, 121.51058823529411, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(327.54146341463417, 147.09176470588235, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(331.1707317073171, 159.88235294117646, 'gini = 0.0\nsamples = 16\nvalue = [16, 0]')

```



```
In [40]: dt2.score(df_train3, df_train_y)
```

```
Out[40]: 0.9919678714859438
```

```
In [41]: dt2.score(df_val3, df_val_y)
```

```
Out[41]: 0.768
```

- The model is overfitting for default setting.


3-2-b. Improve a decision tree model by tuning some hyperparameters for terminating the splitting process. Tune the maximum depth of the tree

```
In [42]: depth = []
accuracy = []
for max_depth in range(1,11):
    dt = DecisionTreeClassifier(random_state=0, max_depth=max_depth)
    dt = dt.fit(df_train3, df_train_y)
    print("Max Depth: ", max_depth)
    print("Train Accuracy: ", dt.score(df_train3, df_train_y))
    print("Validation Accuracy: ", dt.score(df_val3, df_val_y))
    print('-'*20)
    depth.append(max_depth)
    accuracy.append(dt.score(df_val3, df_val_y))

plt.plot(depth, accuracy)
plt.title("Accuracy on the validation set for each depth")
plt.xlabel("Max Depth")
plt.ylabel("Validation Accuracy")
```

```
Max Depth: 1
Train Accuracy: 0.7751004016064257
Validation Accuracy: 0.792
-----
Max Depth: 2
Train Accuracy: 0.7811244979919679
Validation Accuracy: 0.856
-----
Max Depth: 3
Train Accuracy: 0.8273092369477911
Validation Accuracy: 0.848
-----
Max Depth: 4
Train Accuracy: 0.8413654618473896
Validation Accuracy: 0.816
-----
Max Depth: 5
Train Accuracy: 0.8534136546184738
Validation Accuracy: 0.824
-----
Max Depth: 6
Train Accuracy: 0.8815261044176707
Validation Accuracy: 0.784
-----
Max Depth: 7
Train Accuracy: 0.8975903614457831
Validation Accuracy: 0.816
-----
Max Depth: 8
Train Accuracy: 0.9156626506024096
Validation Accuracy: 0.816
-----
Max Depth: 9
Train Accuracy: 0.9397590361445783
Validation Accuracy: 0.808
-----
Max Depth: 10
Train Accuracy: 0.9558232931726908
Validation Accuracy: 0.776
-----
```

```
Out[42]: Text(0, 0.5, 'Validation Accuracy')
```



- According to the plot, the optimum maximum depth is 2.

3-2-C. Train a decision tree classifier using the optimum setting of maximum depth found in 3.2b on all data from the training and validation set.

```
In [43]: # Combine training and validation set
df_train_val = pd.concat([df_train3, df_val3])
df_train_val_y = pd.concat([df_train_y, df_val_y])

# Set Decision tree classifier
dt = DecisionTreeClassifier(random_state=1, max_depth = 2)

# Train model
dt.fit(df_train_val, df_train_val_y)
```

Out[43]: DecisionTreeClassifier(max_depth=2, random_state=1)

```
In [44]: # Predict
y_pred = dt.predict(df_test3)
y_pred

# Create data frame for Prediction
df_pred = pd.DataFrame(y_pred)
df_pred

# save the predictions
df_pred.to_csv('F0034WQ_dt_prediction.csv', index=False, header = False)
```

4 Random Forest Classification

4-1. Initialize a random forest classifier.

```
In [45]: from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(max_depth=2, random_state=0)
rf.fit(df_train3, df_train_y)

# Accuracy on Validation set in default set
print("Validation Accuracy: ", rf.score(df_val3, df_val_y))
```

Validation Accuracy: 0.816

4-2. Optimize your classifier by tuning two hyperparameters to maximize the accuracy of your model:

```
In [46]: from sklearn.model_selection import GridSearchCV
```

```
# Create the grid search hyperparameter
grid_hyper = {'max_depth': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
              'n_estimators': [50, 100, 150, 200, 250]}

# classifier
rf = RandomForestClassifier(random_state = 2)

# search different combinations.
rf_GRCV = GridSearchCV(estimator = rf, param_grid = grid_hyper, cv = 10)

# Fit the model
rf_GRCV.fit(df_train_val, df_train_val_y)

# best hyperparameters
print("Optimum parameters: ", rf_GRCV.best_params_)
```

Optimum parameters: {'max_depth': 7, 'n_estimators': 150}

4-3. Report and interpret the feature importance of each feature by your random forest model.

```
In [47]: # classifier with best hyperparameters
rf = RandomForestClassifier(random_state = 2, max_depth = 7, n_estimators = 150)
rf.fit(df_train_val, df_train_val_y)
```

```
Out[47]: RandomForestClassifier(max_depth=7, n_estimators=150, random_state=2)
```

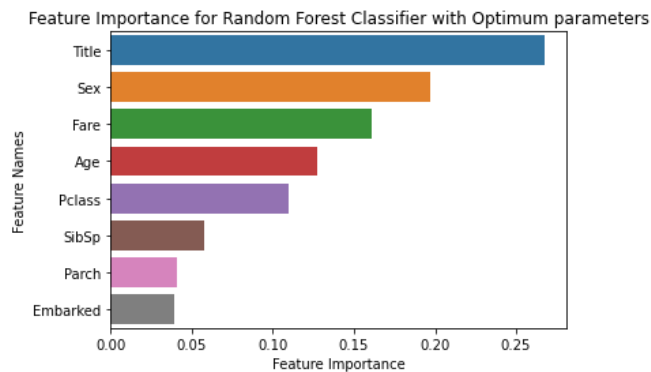
```
In [48]: # Generate feature importance and feature names
feature_importance = np.array(rf.feature_importances_)
feature_names = np.array(df_train_val.columns)

# Generate a DataFrame for feature importance
data={'feature_names':feature_names,'feature_importance':feature_importance}
feature_importance_df = pd.DataFrame(data)

# Sort the DataFrame according to feature importance
feature_importance_df.sort_values(by=['feature_importance'], ascending=False, inplace=True)

# Plot
import seaborn as sns
sns.barplot(x = feature_importance_df['feature_importance'], y = feature_importance_df['feature_names'])
plt.title('Feature Importance for Random Forest Classifier with Optimum parameters')
plt.xlabel('Feature Importance')
plt.ylabel('Feature Names')
```

```
Out[48]: Text(0, 0.5, 'Feature Names')
```



- According to the above plot, we can see the most important feature in my model is "Title".

4-4. Predict the survivals on testing data using the classifier with the optimum parameters from 4.2.

```
In [49]: # Predict
y_pred = rf.predict(df_test3)
y_pred

# Create data frame for Prediction
df_pred = pd.DataFrame(y_pred)
df_pred

# save the predictions
df_pred.to_csv('F0034WQ_rf_prediction.csv', index=False, header = False)
```

5. SVM classification

```
In [50]: # Pre-processed data
from sklearn.preprocessing import StandardScaler
SS = StandardScaler()

df_train3_SS = SS.fit_transform(df_train3)
df_val3_SS = SS.transform(df_val3)
df_test3_SS = SS.transform(df_test3)
```

5-1. Build SVM classifiers with the following three kernels: linear, polynomial with degree = 2, and radial basis function kernel. Train the classifiers and calculate the performance of each classifier on the validation set.

```
In [51]: from sklearn.metrics import accuracy_score, log_loss
from sklearn.svm import SVC

svm_linear = SVC(kernel = 'linear')
svm_poly = SVC(kernel = 'poly', degree = 2, probability = True)
svm_rbf = SVC(kernel = 'rbf', probability = True)

# kernels: linear
svm_linear.fit(df_train3_SS, df_train_y)
print("kernels: linear")
print("Validation Accuracy: ", svm_linear.score(df_val3_SS, df_val_y))
print('-'*20)

# kernels: polynomial with degree = 2
svm_poly.fit(df_train3_SS, df_train_y)
print("kernels: polynomial with degree = 2")
print("Validation Accuracy: ", svm_poly.score(df_val3_SS, df_val_y))
print('Loss of Validation Set:', log_loss(df_val_y, svm_poly.predict_proba(df_val3_SS)) )
print('-'*20)

# kernels: radial basis function
svm_rbf.fit(df_train3_SS, df_train_y)
print("kernels: radial basis function")
print("Validation Accuracy: ", svm_rbf.score(df_val3_SS, df_val_y))
print('Loss of Validation Set:', log_loss(df_val_y, svm_rbf.predict_proba(df_val3_SS)) )
print('-'*20)

kernels: linear
Validation Accuracy:  0.84
-----
kernels: polynomial with degree = 2
Validation Accuracy:  0.856
Loss of Validation Set: 0.4049651318151844
-----
kernels: radial basis function
Validation Accuracy:  0.856
```


Loss of Validation Set: 0.3780804868895341

5-2. Which kernel gives the best result? Use your classifier with the this kernel to predict the survival outcomes in the testing set.

- Although classifier with kernel = polynomial and classifier with kernel = radial basis function return same accuracy, classifier with kernel = radial basis function returns smaller Loss. As a result, classifier with kernel = radial basis function gives the best result.

```
In [52]: # Predict
y_pred = svm_rbf.predict(df_test3_SS)
y_pred

# Create data frame for Prediction
df_pred = pd.DataFrame(y_pred)
df_pred

# save the predictions
df_pred.to_csv('F0034WQ_svm_prediction.csv', index=False, header = False)
```

In []: